

File Systems

Agenda

- Basics
 - File system implementation
 - File allocation algorithms
- Introduction to UNIX file system

File system implementation

- A disk usually contains multiple partitions
 - These are created when installing the OS
 - Each partition can either be a raw partition or hold a file system
 - Raw partitions (containing unformatted characters) are often used to create swap space in UNIX
 - A file system partition contains four types of information
 - A boot control block or boot block or boot sector: the first block of the partition
 - A superblock or volume control block or master file table: holds various metadata of the file system
 - A directory structure
 - The data blocks containing the file data

File allocation

- File allocation is done through a mapping from logical file blocks to physical disk sectors
 - Contiguous or linear allocation: logical file blocks are laid out contiguously on the disk
 - Directory entry contains the block address of the first block and the length of the file (block address vs. disk address?)
 - Offers easy random access, but suffers from external fragmentation due to size over-estimate
 - Growing a file beyond its allocated space is impossible
 - Linked allocation: each file block has a pointer to the next block; no external fragmentation
 - Directory entry contains the addresses of the first and the last blocks
 - File allocation table (FAT) of MS-DOS uses a linked allocation with a table to improve random access latency

File allocation

- Mapping from file blocks to disk sectors
 - Indexed allocation: each file has its own index block, which stores the list of block addresses corresponding to the file blocks
 - Directory entry stores the index block number
 - Fast random access and no external fragmentation
 - Small files waste a lot of space in the index block
 - Large files may need a linked list of multiple index blocks with the directory entry storing the head index block id
 - For large files, instead of having a linked list of index blocks, it is possible to have multiple levels of index blocks
 - UNIX uses a multi-level index block scheme; each file has an inode (index node) storing ten direct data block addresses, one single-indirect pointer, one double-indirect pointer, and one triple-indirect pointer

UNIX file system

- Every file has a unique inode
 - Contains information about owner, group, permissions, time of last access, size, data block addresses, and few more pieces of information
 - A directory is treated just like a file where the data blocks hold the file name to inode map for each file under that directory
 - UNIX inodes may reside in disk or in the in-memory inode cache
 - The disk inodes of a partition are arranged in an array of pre-defined size
 - This array is placed between the superblock and the data blocks
 - The size of the array determines the maximum number of files in the file system hosted by the partition

UNIX file system

- When an inode is read from the disk, it is copied into a free inode in the in-memory inode cache
 - Inode cache is a hash table
 - The hash key is a function of inode disk partition number and the inode number
 - A free list chains up all free inodes in the inode cache