

CS330: Operating Systems

Quiz#4

Name:

Roll No.:

1. Consider the following function. The semaphores a, b, c, d are all initialized to one.

```
sem_t a, b, c, d;
void f (sem_t *s1, sem_t *s2, sem_t *s3, sem_t *s4)
{
    sem_wait(s1); sem_wait(s2); sem_wait(s3); sem_wait(s4);
    Critical section
    sem_post(s1); sem_post(s2); sem_post(s3); sem_post(s4);
}
```

Two threads call the function `f` concurrently. One of these threads invokes `f` as `f(&c, &b, &a, &d)`. Write down all possible correct invocations of function `f` by the other thread. (2 points)

Solution: All invocations that share a common prefix with `f(&c, &b, &a, &d)` will be correct. So, these are `f(&c, &b, &a, &d)`, `f(&c, &b, &d, &a)`, `f(&c, &a, &b, &d)`, `f(&c, &a, &d, &b)`, `f(&c, &d, &a, &b)`, `f(&c, &d, &b, &a)`.

Grading policy: 0.5 mark for writing any one of the six correct invocations. One mark for writing at least two correct invocations. 2 marks for writing all six. Zero mark for writing anything wrong.

2. The following program segment is executed by ten threads. The initial value of the shared variable `x` is zero. The semaphore `s` is initialized to eight. What are the possible values of `x` after all threads complete execution? (3 points)

```
sem_wait(&s);
x++;
sem_post(&s);
```

Solution: At any point in time, there can be at most eight threads executing `x++` concurrently. Thus during the execution of the first batch of eight threads, the value of `x` can vary from 1 to 8. Therefore, when the remaining two threads execute, the value of `x` can go to 9 or 10. Also, it is possible that the thread, which read the value of `x` as zero during the execution of the first batch of threads, finishes last and writes 1 to `x`. Thus, the possible values of `x` are 1, 2, 3, 4, 5, 6, 7, 8, 9, 10. To see how `x` can attain a final value of k for $1 \leq k \leq 10$, consider a thread T that reads `x` after $k - 1$ threads have read, incremented, and updated `x` one by one maintaining mutual exclusion. It is always possible to construct such a schedule as long as the initial value of the semaphore is positive. So, the value of `x` read by T is $k - 1$. Now, you need to construct a schedule where T is context switched right after reading `x` and is scheduled again after all other threads have completed execution. Construction of such a schedule requires the semaphore value to be at least two. So, T is the last thread to update `x`. Therefore, the final value of `x` will be k .

Grading policy: One mark less for each missing possible value. Thus, if you miss three or more possible values of x , you get zero.

3. Consider two semaphores s_1 and s_2 . The variable `count` is shared between two threads and initialized to x . The following program segments are executed by two threads. If the final value of `count` is always $x-1$ after two threads have completed execution, what are the initial values of the semaphores s_1 and s_2 ? **(3 points)**

Thread A	Thread B
-----	-----
<code>sem_wait(&s1);</code>	<code>sem_wait(&s2);</code>
<code>if (count == x) count--;</code>	<code>if (count == x) count -= 2;</code>
<code>else count++;</code>	<code>else count += 2;</code>
<code>sem_post(&s2);</code>	<code>sem_post(&s1);</code>

Solution: For `count` to have a final value of $x-1$, Thread B must execute first and then Thread A executes. So, s_2 should have a positive initial value and s_1 should have a zero initial value.

Grading policy: 1.5 marks for each correct answer. Zero mark if you write multiple possible initial values for s_1 .

4. Consider the following program segments executed by two threads. What possible values can Thread B print? **(2 points)**

<code>pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;</code>	
<code>pthread_cond_t cv = PTHREAD_COND_INITIALIZER;</code>	
<code>int x = 0;</code>	
<code>int y = 0;</code>	
Thread A	Thread B
-----	-----
<code>y = 2;</code>	<code>while (x == 0) {</code>
<code>pthread_mutex_lock(&lock);</code>	<code>pthread_mutex_lock(&lock);</code>
<code>x = 1;</code>	<code>pthread_cond_wait(&cv, &lock);</code>
<code>pthread_cond_signal(&cv);</code>	<code>pthread_mutex_unlock(&lock);</code>
<code>pthread_mutex_unlock(&lock);</code>	<code>}</code>
	<code>printf ("%d\n", y);</code>

Solution: If Thread A executes first or starts executing while Thread B is sleeping on the condition variable, the printed value will be 2. If Thread B executes first, checks the while loop condition, and then switches context before acquiring the lock, Thread A may execute and complete. Now, when Thread B is scheduled, it will keep sleeping on the condition variable forever. Thus, the second possibility is a deadlock with Thread B printing nothing.

To avoid the second possibility, Thread B should acquire the lock before checking the while loop condition. By not doing so, it has introduced a data race between the read of x and the write of x .

Grading policy: One mark for clearly mentioning each possibility.